


LINEAR REGRESSION

Chih-Chung Hsu (許志仲)
Assistant Professor
ACVLab, Institute of Data Science
National Cheng Kung University
<https://cchsui.info>



Supervised learning

- Given a set of **training examples**: $x_i = \langle x_{i1}, x_{i2}, x_{i3}, \dots, x_{in}, y_i \rangle$

x_{ij} is the j^{th} feature of the i^{th} example

y_i is the desired **output** (or **target**) for the i^{th} example.

X_j denotes the j^{th} feature.

- We want to learn a function $f: X_1 \times X_2 \times \dots \times X_n \rightarrow Y$
which maps the input variables onto the output domain.

| tumor size | texture | perimeter | ... | outcome | time |
|------------|---------|-----------|-----|---------|------|
| 18.02 | 27.6 | 117.5 | | N | 31 |
| 17.99 | 10.38 | 122.8 | | N | 61 |
| 20.29 | 14.34 | 135.1 | | R | 27 |
| ... | | | | | |

Supervised learning

- Given a dataset $X \times Y$, find a function: $f: X \rightarrow Y$ such that $f(\mathbf{x})$ is a good predictor for the value of y .
- Formally, f is called the **hypothesis**.
- Output Y can have many types:
 - If $Y = \mathbb{R}$, this problem is called **regression**.
 - If Y is a finite discrete set, the problem is called **classification**.
 - If Y has 2 elements, the problem is called **binary classification**.

Prediction problems

- The problem of predicting tumour recurrence is called:
classification
- The problem of predicting the time of recurrence is called:
regression
- Treat them as two separate supervised learning problems.


| tumor size | texture | perimeter | ... | outcome | time |
|------------|---------|-----------|-----|---------|------|
| 18.02 | 27.6 | 117.5 | | N | 31 |
| 17.99 | 10.38 | 122.8 | | N | 61 |
| 20.29 | 14.34 | 135.1 | | R | 27 |
| ... | | | | | |

Variable types

- **Quantitative**, often real number measurements.
 - Assumes that similar measurements are similar in nature.
- **Qualitative**, from a set (categorical, discrete).
 - E.g. {Spam, Not-spam}
- **Ordinal**, also from a discrete set, without metric relation, but that allows ranking.
 - E.g. {first, second, third}

A regression problem

- What hypothesis class should we pick?



| <i>Observe</i> | <i>Predict</i> |
|----------------|----------------|
| <u>x</u> | <u>y</u> |
| 0.86 | 2.49 |
| 0.09 | 0.83 |
| -0.85 | -0.25 |
| 0.87 | 3.10 |
| -0.44 | 0.87 |
| -0.43 | 0.02 |
| -1.1 | -0.12 |
| 0.40 | 1.81 |
| -0.96 | -0.83 |
| 0.17 | 0.43 |

Linear hypothesis

- Suppose Y is a **linear function** of X :

$$\begin{aligned}f_w(\mathbf{X}) &= w_0 + w_1 x_1 + \dots + w_m x_m \\&= w_0 + \sum_{j=1:m} w_j x_j\end{aligned}$$

- The w_j are called **parameters** or **weights**.
- To simplify notation, we add an attribute $x_0=1$ to the m other attributes (also called **bias term** or **intercept**).

How should we pick the *weights*?

Least-squares solution method


- The linear regression problem: $f_w(X) = w_0 + \sum_{j=1:m} w_j x_j$
where m = the dimension of observation space, i.e. number of features.

- **Goal:** Find the **best** linear model given the data.
- Many different possible **evaluation** criteria!
- Most common choice is to find the w that minimizes:

$$Err(w) = \sum_{i=1:n} (y_i - w^T x_i)^2$$

(A note on notation: Here w and x are column vectors of size $m+1$.)

Least-squares solution for $X \in \mathbb{R}^2$



Least-squares solution method

- Re-write in matrix notation: $f_w(X) = Xw$

$$Err(w) = (Y - Xw)^T(Y - Xw)$$

where X is the $n \times m$ matrix of input data,
 Y is the $n \times 1$ vector of output data,
 w is the $m \times 1$ vector of weights.

- To minimize, take the derivative w.r.t. w :

$$\partial Err(w) / \partial w = -2 X^T (Y - Xw)$$

- You get a system of m equations with m unknowns.


- Set these equations to 0: $X^T (Y - Xw) = 0$

- Remember that derivative has to be 0 at a minimum of $Err(w)$

Least-squares solution method

- We want to solve for \mathbf{w} : $X^T (Y - X\mathbf{w}) = 0$
- Try a little algebra: $X^T Y = X^T X \mathbf{w}$
 $\hat{\mathbf{w}} = (X^T X)^{-1} X^T Y$
($\hat{\mathbf{w}}$ denotes the estimated weights)
- Train set predictions: $\hat{Y} = X\hat{\mathbf{w}} = X (X^T X)^{-1} X^T Y$
- Predict new data $X' \rightarrow Y'$: $Y' = X'\hat{\mathbf{w}} = X' (X^T X)^{-1} X^T Y$

Example of linear regression



| x | y |
|-------|-------|
| 0.86 | 2.49 |
| 0.09 | 0.83 |
| -0.85 | -0.25 |
| 0.87 | 3.10 |
| -0.44 | 0.87 |
| -0.43 | 0.02 |
| -1.10 | -0.12 |
| 0.40 | 1.81 |
| -0.96 | -0.83 |
| 0.17 | 0.43 |

What is a plausible estimate of w ?

Try it!

Data matrices

$$\begin{aligned} X^T X &= \begin{bmatrix} 0.86 & 0.09 & -0.85 & 0.87 & -0.44 & -0.43 & -1.10 & 0.40 & -0.96 & 0.17 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \end{bmatrix} \times \begin{bmatrix} 0.86 & 1 \\ 0.09 & 1 \\ -0.85 & 1 \\ 0.87 & 1 \\ -0.44 & 1 \\ -0.43 & 1 \\ -1.10 & 1 \\ 0.40 & 1 \\ -0.96 & 1 \\ 0.17 & 1 \end{bmatrix} \\ &= \begin{bmatrix} 4.95 & -1.39 \\ -1.39 & 10 \end{bmatrix} \end{aligned}$$

Data matrices

$$X^T Y =$$
$$\begin{bmatrix} 0.86 & 0.09 & -0.85 & 0.87 & -0.44 & -0.43 & -1.10 & 0.40 & -0.96 & 0.17 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \end{bmatrix} \times \begin{bmatrix} 2.49 \\ 0.83 \\ -0.25 \\ 3.10 \\ 0.87 \\ 0.02 \\ -0.12 \\ 1.81 \\ -0.83 \\ 0.43 \end{bmatrix}$$
$$= \begin{bmatrix} 6.49 \\ 8.34 \end{bmatrix}$$

Solving the problem


$$\mathbf{w} = (X^T X)^{-1} X^T Y = \begin{bmatrix} 4.95 & -1.39 \\ -1.39 & 10 \end{bmatrix}^{-1} \begin{bmatrix} 6.49 \\ 8.34 \end{bmatrix} = \begin{bmatrix} 1.60 \\ 1.05 \end{bmatrix}$$

So the best fit line is $y = 1.60x + 1.05$.

Solving the problem

$$\mathbf{w} = (X^T X)^{-1} X^T Y = \begin{bmatrix} 4.95 & -1.39 \\ -1.39 & 10 \end{bmatrix}^{-1} \begin{bmatrix} 6.49 \\ 8.34 \end{bmatrix} = \begin{bmatrix} 1.60 \\ 1.05 \end{bmatrix}$$

So the best fit line is $y = 1.60x + 1.05$.



Interpreting the solution

- Linear fit for a prostate cancer dataset
 - Features $X = \{\text{lcavol}, \text{lweight}, \text{age}, \text{lbph}, \text{svi}, \text{lcp}, \text{gleason}, \text{pgg45}\}$
 - Output y = level of PSA (an enzyme which is elevated with cancer).
 - High coefficient weight (in absolute value) = important for prediction.

| Variable | Description |
|----------|-------------------------------------|
| lcavol | (log) Cancer Volume |
| lweight | (log) Weight |
| age | Patient age |
| lbph | (log) Venning Prostatic Hyperplasia |
| svi | Seminal Vesicle Invasion |
| lcp | (log) Capsular Penetration |
| gleason | Gleason score |
| pgg45 | Percent of Gleason score 4 or 5 |
| lpsa | (log) Prostate Specific Antigen |
| train | Label for test / training split |

| Term | Coefficient | Std. Error |
|-----------|--------------|------------|
| Intercept | $w_0 = 2.46$ | 0.09 |
| lcavol | 0.68 | 0.13 |
| lweight | 0.26 | 0.10 |
| age | -0.14 | 0.10 |
| lbph | 0.21 | 0.10 |
| svi | 0.31 | 0.12 |
| lcp | -0.29 | 0.15 |
| gleason | -0.02 | 0.15 |
| pgg45 | 0.27 | 0.15 |

Interpreting the solution

- Caveat: data should be in same range
- If we change unit for age from years to months, we expect the optimal weight to be 12x as low (so predictions don't change)
- Doesn't mean age became 12x less relevant!
- Can **normalize** data to make range similar
 - E.g. subtract average and divide by standard deviation
- More principled approach in next lecture

Computational cost of linear regression

- What operations are necessary?

Computational cost of linear regression

- What operations are necessary?
 - Overall: 1 matrix inversion + 3 matrix multiplications
 - $X^T X$ (other matrix multiplications require fewer operations.)
 - X^T is $m \times n$ and X is $n \times m$, so we need nm^2 operations.
 - $(X^T X)^{-1}$
 - $X^T X$ is $m \times m$, so we need m^3 operations.

Computational cost of linear regression

- What operations are necessary?
 - Overall: 1 matrix inversion + 3 matrix multiplications
 - $X^T X$ (other matrix multiplications require fewer operations.)
 - X^T is $m \times n$ and X is $n \times m$, so we need nm^2 operations.
 - $(X^T X)^{-1}$
 - $X^T X$ is $m \times m$, so we need m^3 operations.
- We can do linear regression in polynomial time, but handling large datasets (many examples, many features) can be problematic.

An alternative for minimizing mean-squared error (MSE)

- Recall the least-square solution: $\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$
- What if \mathbf{X} is too big to compute this explicitly (e.g. $m \sim 10^6$)?

An alternative for minimizing mean-squared error (MSE)

- Recall the least-square solution: $\hat{\mathbf{w}} = (X^T X)^{-1} X^T Y$
- What if X is too big to compute this explicitly (e.g. $m \sim 10^6$)?


- Go back to the gradient step: $Err(\mathbf{w}) = (Y - X\mathbf{w})^T (Y - X\mathbf{w})$

$$\partial Err(\mathbf{w}) / \partial \mathbf{w} = -2 X^T (Y - X\mathbf{w})$$

$$\partial Err(\mathbf{w}) / \partial \mathbf{w} = 2(X^T X \mathbf{w} - X^T Y)$$

Gradient-descent solution for MSE

- Consider the error function:



- The gradient of the error is a vector indicating the direction to the minimum point.
- Instead of directly finding that minimum (using the closed-form equation), we can take small steps towards the minimum.

Gradient-descent solution for MSE

- We want to produce a sequence of weight solutions, $\mathbf{w}_0, \mathbf{w}_1, \mathbf{w}_2 \dots$, such that: $Err(\mathbf{w}_0) > Err(\mathbf{w}_1) > Err(\mathbf{w}_2) > \dots$

Gradient-descent solution for MSE

- We want to produce a sequence of weight solutions, $\mathbf{w}_0, \mathbf{w}_1, \mathbf{w}_2, \dots$, such that: $Err(\mathbf{w}_0) > Err(\mathbf{w}_1) > Err(\mathbf{w}_2) > \dots$

- The algorithm:

Given an initial weight vector \mathbf{w}_0 ,

Do for $k=1, 2, \dots$

$$\mathbf{w}_{k+1} = \mathbf{w}_k - \alpha_k \partial Err(\mathbf{w}_k) / \partial \mathbf{w}_k$$

End when $|\mathbf{w}_{k+1} - \mathbf{w}_k| < \epsilon$

- Parameter $\alpha_k > 0$ is the step-size (or learning rate) for iteration k .

Convergence

- Convergence depends in part on the α_k .
- **If steps are too large**: the \mathbf{w}_k may oscillate forever.
 - This suggests that $\alpha_k \rightarrow 0$ as $k \rightarrow \infty$.
- **If steps are too small**: the \mathbf{w}_k may not move far enough to reach a local minimum.

Robbins-Monroe conditions

- The α_k are a Robbins-Monroe sequence if:

$$\sum_{k=0:\infty} \alpha_k = \infty$$

$$\sum_{k=0:\infty} \alpha_k^2 < \infty$$

- These conditions are sufficient to ensure convergence of the \mathbf{w}_k to a **local minimum** of the error function.

Robbins-Monroe conditions

- The α_k are a Robbins-Monroe sequence if:

$$\sum_{k=0:\infty} \alpha_k = \infty$$

$$\sum_{k=0:\infty} \alpha_k^2 < \infty$$

- These conditions are sufficient to ensure convergence of the \mathbf{w}_k to a **local minimum** of the error function.

E.g. $\alpha_k = 1 / (k + 1)$ (averaging)


E.g. $\alpha_k = 1/2$ for $k = 1, \dots, T$

$\alpha_k = 1/2^2$ for $k = T+1, \dots, (T+1)+2T$

etc.

Local minima


- Convergence is **NOT** to a global minimum, only to local minimum.



- The blue line represents the **error function**. There is no guarantee regarding the amount of error of the weight vector found by gradient descent, compared to the globally optimal solution.

Local minima


- Convergence is **NOT** to a global minimum, only to local minimum.



- For linear function approximations using Least-Mean Squares (LMS) error, this is not an issue: **only ONE global minimum!**
 - Local minima affects many other function approximators.

Local minima

- Convergence is **NOT** to a global minimum, only to local minimum.



- For linear function approximations using Least-Mean Squares (LMS) error, this is not an issue: **only ONE global minimum!**
 - Local minima affects many other function approximators.
- Repeated random restarts can help (in all cases of gradient search).

Basic least-squares solution method

- Recall the least-square solution: $\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$
- Assuming for now that \mathbf{X} is reasonably small so computation and memory are not a problem. Can we always evaluate this?

Basic least-squares solution method

- Recall the least-square solution: $\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$
- Assuming for now that \mathbf{X} is reasonably small so computation and memory are not a problem. Can we always evaluate this?
- To have a unique solution, we need $\mathbf{X}^T \mathbf{X}$ to be nonsingular.
That means \mathbf{X} must have full column rank (i.e. no features can be expressed using other features.)

Exercise: What if \mathbf{X} does not have full column rank? When would this happen? Design an example. Try to solve it.

Dealing with difficult cases of $(X^T X)^{-1}$

- **Case #1:** The weights are not uniquely defined.

Solution: Re-code or drop some redundant columns of X .


- **Case #2:** The number of features/weights (m) exceeds the number of training examples (n).

Solution: Reduce the number of features using various techniques (to be studied later.)

Predicting recurrence time from tumor size

This function looks complicated, and a linear hypothesis does not seem very good.

What should we do?




Predicting recurrence time from tumor size

This function looks complicated, and a linear hypothesis does not seem very good.

What should we do?

- *Pick a better function?*
- *Use more features?*
- *Get more data?*




Input variables for linear regression

- Original quantitative variables X_1, \dots, X_m
- Transformations of variables, e.g. $X_{m+1} = \log(X_i)$
- Basis expansions, e.g. $X_{m+1} = X_i^2, X_{m+2} = X_i^3, \dots$
- Interaction terms, e.g. $X_{m+1} = X_i X_j$
- Numeric coding of qualitative variables, e.g. $X_{m+1} = 1$ if X_i is true and 0 otherwise.

In all cases, we can add X_{m+1}, \dots, X_{m+k} to the list of original variables and perform the linear regression.

Example of linear regression with polynomial terms

$$f_w(x) = w_0 + w_1 x + w_2 x^2$$




$$X = \begin{bmatrix} x^2 & x & 1 \\ 0.75 & 0.86 & 1 \\ 0.01 & 0.09 & 1 \\ 0.73 & -0.85 & 1 \\ 0.76 & 0.87 & 1 \\ 0.19 & -0.44 & 1 \\ 0.18 & -0.43 & 1 \\ 1.22 & -1.10 & 1 \\ 0.16 & 0.40 & 1 \\ 0.93 & -0.96 & 1 \\ 0.03 & 0.17 & 1 \end{bmatrix} \quad Y = \begin{bmatrix} 2.49 \\ 0.83 \\ -0.25 \\ 3.10 \\ 0.87 \\ 0.02 \\ -0.12 \\ 1.81 \\ -0.83 \\ 0.43 \end{bmatrix}$$

Solving the problem

$$\mathbf{w} = (X^T X)^{-1} X^T Y = \begin{bmatrix} 4.11 & -1.64 & 4.95 \\ -1.64 & 4.95 & -1.39 \\ 4.95 & -1.39 & 10 \end{bmatrix}^{-1} \begin{bmatrix} 3.60 \\ 6.49 \\ 8.34 \end{bmatrix} = \begin{bmatrix} 0.68 \\ 1.74 \\ 0.73 \end{bmatrix}$$

So the best order-2 polynomial is $y = 0.68x^2 + 1.74x + 0.73$.



Compared to $y = 1.6x + 1.05$
for the order-1 polynomial.

Input variables for linear regression


How to choose input variables?

- Propose different strategies, then perform model selection using cross validation (more details later)
- Add many transformation to the set of features, then perform feature selection or dimension reduction (more details later)
- Use problem specific insights:
 - Say, predict displacement of falling object as function of time
 - From physics, know that $s=gt^2$
 - In that case, use squared transformation of t (input variable is t^2)

What you should know

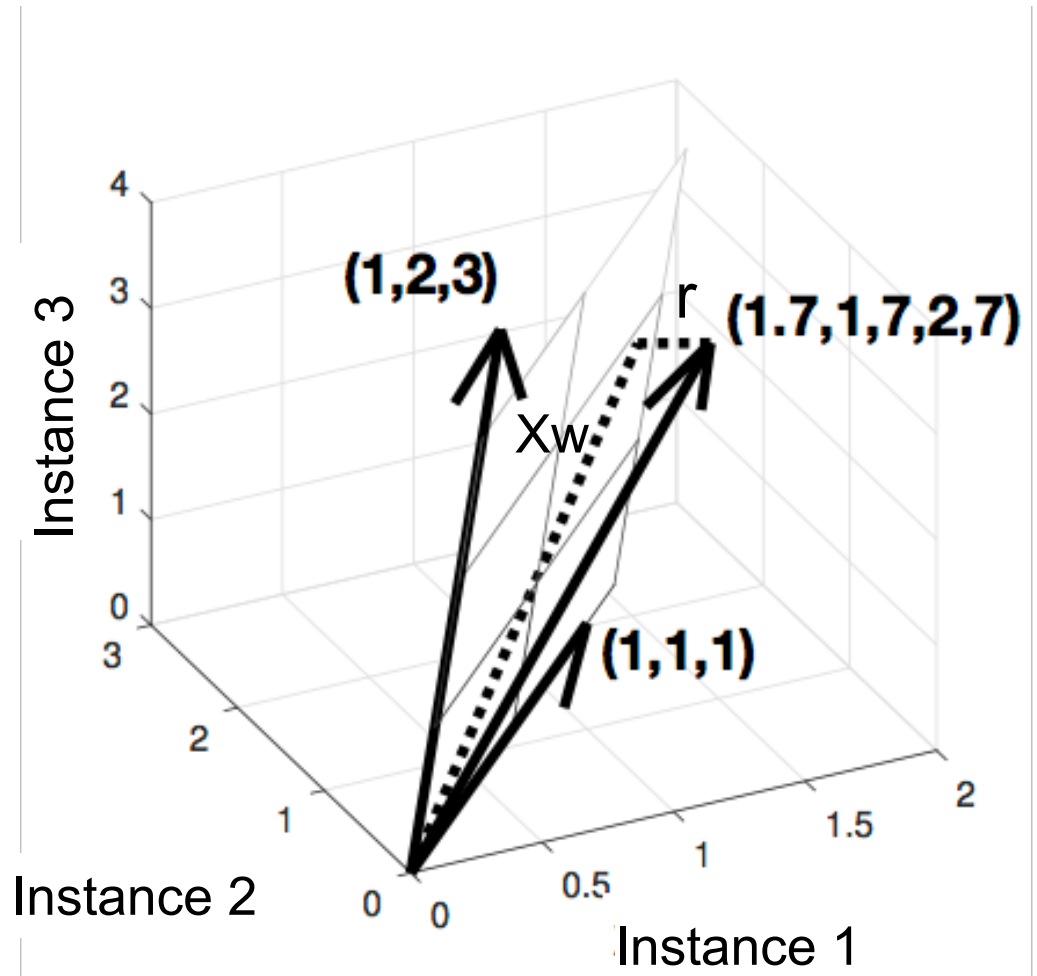
- Definition and characteristics of a supervised learning problem.
- Linear regression (hypothesis class, cost function).
- Closed-form least-squares solution method (algorithm, computational complexity, stability issues).
- Gradient descent method (algorithm, properties).

Weight space view



Instance space view (Geometric view)


$$\begin{matrix} & X & & y \\ \begin{bmatrix} 1 & 1 \\ 1 & 2 \\ 1 & 3 \end{bmatrix} & \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} & \approx & \begin{bmatrix} 1.7 \\ 1.7 \\ 2.7 \end{bmatrix} \end{matrix}$$



Predicting recurrence time from tumor size

This function looks complicated, and a linear hypothesis does not seem very good.

What should we do?




Predicting recurrence time from tumor size

This function looks complicated, and a linear hypothesis does not seem very good.

What should we do?

- *Pick a better function?*
- *Use more features?*
- *Get more data?*




Input variables for linear regression

- Original quantitative variables X_1, \dots, X_m
- Transformations of variables, e.g. $X_{m+1} = \log(X_i)$
- Basis expansions, e.g. $X_{m+1} = X_i^2, X_{m+2} = X_i^3, \dots$
- Interaction terms, e.g. $X_{m+1} = X_i X_j$
- Numeric coding of qualitative variables, e.g. $X_{m+1} = 1$ if X_i is true and 0 otherwise.

In all cases, we can add X_{m+1}, \dots, X_{m+k} to the list of original variables and perform the linear regression.

Example of linear regression with polynomial terms

$$f_w(x) = w_0 + w_1 x + w_2 x^2$$




$$X = \begin{bmatrix} x^2 & x & 1 \\ 0.75 & 0.86 & 1 \\ 0.01 & 0.09 & 1 \\ 0.73 & -0.85 & 1 \\ 0.76 & 0.87 & 1 \\ 0.19 & -0.44 & 1 \\ 0.18 & -0.43 & 1 \\ 1.22 & -1.10 & 1 \\ 0.16 & 0.40 & 1 \\ 0.93 & -0.96 & 1 \\ 0.03 & 0.17 & 1 \end{bmatrix} \quad Y = \begin{bmatrix} 2.49 \\ 0.83 \\ -0.25 \\ 3.10 \\ 0.87 \\ 0.02 \\ -0.12 \\ 1.81 \\ -0.83 \\ 0.43 \end{bmatrix}$$

Solving the problem


$$\mathbf{w} = (X^T X)^{-1} X^T Y = \begin{bmatrix} 4.11 & -1.64 & 4.95 \\ -1.64 & 4.95 & -1.39 \\ 4.95 & -1.39 & 10 \end{bmatrix}^{-1} \begin{bmatrix} 3.60 \\ 6.49 \\ 8.34 \end{bmatrix} = \begin{bmatrix} 0.68 \\ 1.74 \\ 0.73 \end{bmatrix}$$

So the best order-2 polynomial is $y = 0.68x^2 + 1.74x + 0.73$.




Compared to $y = 1.6x + 1.05$
for the order-1 polynomial.


Order-3 fit: Is this better?




Order-4 fit




Order-5 fit




Order-6 fit




Order-7 fit



Order-8 fit



Order-9 fit



This is overfitting!

This is overfitting!


- We can find a hypothesis that explains perfectly the training data, but **does not generalize** well to new data.
- **In this example:** we have a lot of parameters (weights), so the hypothesis matches the data points exactly,
but is wild everywhere else.
- A **very important** problem in machine learning.

Overfitting


- Every hypothesis has a **true error** measured on **all possible data items** we could ever encounter (e.g. $f_w(\mathbf{x}_i) - y_i$).
- Since we don't have all possible data, in order to decide what is a good hypothesis, we **measure error over the training set**.
- Formally: Suppose we compare hypotheses f_1 and f_2 .
 - Assume f_1 has lower error on the training set.
 - If f_2 has lower true error, then our algorithm is overfitting.

Overfitting


- Which hypothesis has the lowest **true** error?




$d=1$




$d=2$




$d=3$




$d=4$




$d=5$



$d=6$



$d=7$



$d=8$


Cross-Validation

- Partition your data into a **Training Set** and a **Validation set**.
 - The proportions in each set can vary.
- Use the **Training Set** to find the best hypothesis in the class.
- Use the **Validation Set** to evaluate the true prediction error.
 - Compare across different hypothesis classes (different order polynomials.)

| | | | | | | |
|------------------|-------|------|-------|---|-------|-------|
| <i>Train:</i> | $X =$ | 0.75 | 0.86 | 1 | $Y =$ | 2.49 |
| | | 0.01 | 0.09 | 1 | | 0.83 |
| | | 0.73 | -0.85 | 1 | | -0.25 |
| | | 0.76 | 0.87 | 1 | | 3.10 |
| | | 0.19 | -0.44 | 1 | | 0.87 |
| | | 0.18 | -0.43 | 1 | | 0.02 |
| | | 1.22 | -1.10 | 1 | | -0.12 |
| <i>Validate:</i> | | 0.16 | 0.40 | 1 | | 1.81 |
| | | 0.93 | -0.96 | 1 | | -0.83 |
| | | 0.03 | 0.17 | 1 | | 0.43 |

k -fold Cross-Validation


- Consider k partitions of the data (usually of equal size).
- Train with $k-1$ subset, validate on k^{th} subset. Repeat k times.
- Average the prediction error over the k rounds/folds.



Source: <http://stackoverflow.com/questions/31947183/how-to-implement-walk-forward-testing-in-sklearn>

k -fold Cross-Validation

- Consider k partitions of the data (usually of equal size).
- Train with $k-1$ subset, validate on k^{th} subset. Repeat k times.
- Average the prediction error over the k rounds/folds.



Source: <http://stackoverflow.com/questions/31947183/how-to-implement-walk-forward-testing-in-sklearn>

- **Computation time is increased** by factor of k .

Leave-one-out cross-validation

- Let $k = n$, the size of the training set
- For each order- d hypothesis class,
 - Repeat n times:
 - Set aside one instance $\langle x_i, y_i \rangle$ from the training set.
 - Use all other data points to find w (optimization).
 - Measure prediction error on the held-out $\langle x_i, y_i \rangle$.
 - Average the prediction error over all n subsets.
- Choose the d with lowest **estimated true prediction error**.

Estimating true error for $d=1$

Data

| x | y |
|-------|-------|
| 0.86 | 2.49 |
| 0.09 | 0.83 |
| -0.85 | -0.25 |
| 0.87 | 3.10 |
| -0.44 | 0.87 |
| -0.43 | 0.02 |
| -1.1 | -0.12 |
| 0.40 | 1.81 |
| -0.96 | -0.83 |
| 0.17 | 0.43 |

Cross-validation results

| Iter | D_{train} | D_{valid} | Error _{train} | Error _{valid} |
|-------|--------------------------|----------------|------------------------|------------------------|
| 1 | $D - \{(0.86, 2.49)\}$ | (0.86, 2.49) | 0.4928 | 0.0044 |
| 2 | $D - \{(0.09, 0.83)\}$ | (0.09, 0.83) | 0.1995 | 0.1869 |
| 3 | $D - \{(-0.85, -0.25)\}$ | (-0.85, -0.25) | 0.3461 | 0.0053 |
| 4 | $D - \{(0.87, 3.10)\}$ | (0.87, 3.10) | 0.3887 | 0.8681 |
| 5 | $D - \{(-0.44, 0.87)\}$ | (-0.44, 0.87) | 0.2128 | 0.3439 |
| 6 | $D - \{(-0.43, 0.02)\}$ | (-0.43, 0.02) | 0.1996 | 0.1567 |
| 7 | $D - \{(-1.10, -0.12)\}$ | (-1.10, -0.12) | 0.5707 | 0.7205 |
| 8 | $D - \{(0.40, 1.81)\}$ | (0.40, 1.81) | 0.2661 | 0.0203 |
| 9 | $D - \{(-0.96, -0.83)\}$ | (-0.96, -0.83) | 0.3604 | 0.2033 |
| 10 | $D - \{(0.17, 0.43)\}$ | (0.17, 0.43) | 0.2138 | 1.0490 |
| mean: | | | 0.2188 | 0.3558 |

Cross-validation results

| d | Error _{train} | Error _{valid} |
|-----|------------------------|------------------------|
| 1 | 0.2188 | 0.3558 |
| 2 | 0.1504 | 0.3095 |
| 3 | 0.1384 | 0.4764 |
| 4 | 0.1259 | 1.1770 |
| 5 | 0.0742 | 1.2828 |
| 6 | 0.0598 | 1.3896 |
| 7 | 0.0458 | 38.819 |
| 8 | 0.0000 | 6097.5 |

- Optimal choice: $d=2$. Overfitting for $d > 2$.

Evaluation

- We use cross-validation for ***model selection***.
- Available labeled data is split into two parts:
 - **Training set** is used to select a hypothesis f from a class of hypotheses F (e.g. regression of a given degree).
 - **Validation set** is used to compare the best f from each hypothesis class across different classes (e.g. different degree regression).
 - Must be untouched during the process of looking for f within a class F .

Evaluation

- After adapting the weights to minimize the error on the **train set**, the weights could be exploiting particularities in the **train set**:
 - have to use the **validation set** as proxy for true error
- After choosing the hypothesis class to minimize error on the **validation set**, the hypothesis class could be adapted to some particularities in the **validation set**
 - **Validation set** is no longer a good proxy for the true error!

Evaluation

- We use cross-validation for ***model selection***.
- Available labeled data is split into parts:
 - **Training set** is used to select a hypothesis f from a class of hypotheses F (e.g. regression of a given degree).
 - **Validation set** is used to compare the best f from each hypothesis class across different classes (e.g. different degree regression).
 - Must be untouched during the process of looking for f within a class F .
- **Test set**: Ideally, a separate set of (labeled) data is withheld to get a true estimate of the generalization error.
 - Cannot be touched during the process of selecting F(Often the “validation set” is called “test set”, without distinction.)

Validation vs Train error

[From Hastie et al. textbook]






FIGURE 2.11. *Test and training error as a function of model complexity.*

Understanding the error

Given set of examples $\langle X, Y \rangle$. Assume that $y = f(x) + \epsilon$, where ϵ is Gaussian noise with zero mean and std deviation σ .



Understanding the error

- Consider standard linear regression solution:


$$Err(\mathbf{w}) = \sum_{i=1:n} (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

- If we consider only the class of **linear** hypotheses, we have systematic prediction error, called **bias**, whenever the data is generated by a non-linear function.
- Depending on what dataset we observed, we may get different solutions. Thus we can also have error due to this **variance**.
 - This occurs even if data is generated from class of linear functions.

An example (from Tom Dietterich)

- The circles are data points. X is drawn uniformly randomly. Y is generated by the function $y = 2\sin(0.5x) + \epsilon$.


Example: 20 points
 $y = x + 2 \sin(1.5x) + N(0,0.2)$



An example (from Tom Dietterich)

- With different sets of 20 points, we get different lines.

50 fits (20 examples each)



Validation vs Train error

[From Hastie et al. textbook]




FIGURE 2.11. *Test and training error as a function of model complexity.*

Gauss-Markov Theorem

- Main result:

The least-squares estimates of the parameters \mathbf{w} have the **smallest variance** among all linear **unbiased** estimates.

Gauss-Markov Theorem

- Main result:

The least-squares estimates of the parameters \mathbf{w} have the **smallest variance** among all linear **unbiased** estimates.

- Understanding the statement:
 - Real parameters are denoted: \mathbf{w}
 - Estimate of the parameters is denoted: $\hat{\mathbf{w}}$
 - Error of the estimator: $Err(\hat{\mathbf{w}}) = E(\hat{\mathbf{w}} - \mathbf{w})^2 = Var(\hat{\mathbf{w}}) + (E(\hat{\mathbf{w}} - \mathbf{w}))^2$
 - Unbiased estimator means: $E(\hat{\mathbf{w}} - \mathbf{w}) = 0$
 - There may exist an estimator that has lower error, but some bias.

Bias vs Variance

- [Gauss-Markov Theorem says:](#)

The least-squares estimates of the parameters \mathbf{w} have the **smallest variance** among all linear **unbiased** estimates.

- **Insight:** Find **lower variance** solution, at the expense of **some bias**.

Bias vs Variance

- Gauss-Markov Theorem says:

The least-squares estimates of the parameters \mathbf{w} have the **smallest variance** among all linear **unbiased** estimates.

- **Insight:** Find **lower variance** solution, at the expense of **some bias**.
- E.g. fix low-relevance weights to 0

Recall our prostate cancer example

- The Z-score measures the effect of dropping that feature from the linear regression.

$$z_j = \hat{w}_j / \text{sqrt}(\sigma^2 v_j)$$

feature,
where \hat{w}_j is the estimated weight of the j^{th}
and v_j is the j^{th} diagonal element of $(X^T X)^{-1}$

TABLE 3.2. Linear model fit to the prostate cancer data. The Z score is the coefficient divided by its standard error (3.12). Roughly a Z score larger than two in absolute value is significantly nonzero at the $p = 0.05$ level.

| Term | Coefficient | Std. Error | Z Score |
|-----------|-------------|------------|---------|
| Intercept | 2.46 | 0.09 | 27.60 |
| lcavol | 0.68 | 0.13 | 5.37 |
| lweight | 0.26 | 0.10 | 2.75 |
| age | -0.14 | 0.10 | -1.40 |
| lbph | 0.21 | 0.10 | 2.06 |
| svi | 0.31 | 0.12 | 2.47 |
| lcp | -0.29 | 0.15 | -1.87 |
| gleason | -0.02 | 0.15 | -0.15 |
| pgg45 | 0.27 | 0.15 | 1.74 |

[From Hastie et al. textbook]

Subset selection

- **Idea:** Keep only a small set of features with non-zero weights.
- **Goal:** Find lower variance solution, at the expense of some bias.
- There are many different methods for choosing subsets.
(More on this later...)
- Least-squares regression can be used to estimate the weights of the selected features.
- Bias as true model might rely on the discarded features!

Bias vs Variance

- Find **lower variance** solution, at the expense of **some bias**.
- Force some weights to 0
- E.g. Include **penalty for model complexity** in error to reduce overfitting.

$$Err(w) = \sum_{i=1:n} (y_i - \mathbf{w}^T \mathbf{x}_i)^2 + \lambda |model_size|$$

λ is a hyper-parameter that controls penalty size.

Ridge regression (aka L2-regularization)

- Constrains the weights by imposing a penalty on their size:

$$\hat{\mathbf{w}}^{ridge} = \underset{\mathbf{w}}{\operatorname{argmin}} \left\{ \sum_{i=1:n} (y_i - \mathbf{w}^T \mathbf{x}_i)^2 + \lambda \sum_{j=0:m} w_j^2 \right\}$$

where λ can be selected manually, or by cross-validation.

Ridge regression (aka L2-regularization)

- Constrains the weights by imposing a penalty on their size:

$$\hat{\mathbf{w}}^{ridge} = \underset{\mathbf{w}}{\operatorname{argmin}} \{ \sum_{i=1:n} (y_i - \mathbf{w}^T \mathbf{x}_i)^2 + \lambda \sum_{j=0:m} w_j^2 \}$$

where λ can be selected manually, or by cross-validation.

- Do a little algebra to get the solution: $\hat{\mathbf{w}}^{ridge} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{Y}$

Ridge regression (aka L2-regularization)

- Re-write in matrix notation: $f_{\mathbf{w}}(X) = X\mathbf{w}$

$$Err(\mathbf{w}) = (Y - X\mathbf{w})^T(Y - X\mathbf{w}) + \lambda\mathbf{w}^T\mathbf{w}$$

- To minimize, take the derivative w.r.t. \mathbf{w} :

$$\partial Err(\mathbf{w})/\partial \mathbf{w} = -2 X^T (Y - X\mathbf{w}) - 2\lambda\mathbf{w} = \mathbf{0}$$

- Try a little algebra:

Ridge regression (aka L2-regularization)

- Re-write in matrix notation: $f_{\mathbf{w}}(X) = X\mathbf{w}$

$$Err(\mathbf{w}) = (Y - X\mathbf{w})^T(Y - X\mathbf{w}) + \lambda\mathbf{w}^T\mathbf{w}$$

- To minimize, take the derivative w.r.t. \mathbf{w} :

$$\partial Err(\mathbf{w})/\partial\mathbf{w} = -2 X^T (Y - X\mathbf{w}) - 2\lambda\mathbf{w} = \mathbf{0}$$

- Try a little algebra:

$$X^T Y = (X^T X + I\lambda) \mathbf{w}$$

$$\hat{\mathbf{w}} = (X^T X + I\lambda)^{-1} X^T Y$$

Ridge regression (aka L2-regularization)

- Constrains the weights by imposing a penalty on their size:

$$\hat{\mathbf{w}}^{ridge} = \underset{\mathbf{w}}{\operatorname{argmin}} \{ \sum_{i=1:n} (y_i - \mathbf{w}^T \mathbf{x}_i)^2 + \lambda \sum_{j=0:m} w_j^2 \}$$

where λ can be selected manually, or by cross-validation.

- Do a little algebra to get the solution: $\hat{\mathbf{w}}^{ridge} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{Y}$
 - The ridge solution is **not equivariant** under scaling of the data, so typically **need to normalize the inputs** first.
 - Ridge gives a smooth solution, effectively shrinking the weights, but drives few weights to 0.

Lasso regression (aka L1-regularization)

- Constrains the weights by penalizing the absolute value of their size:

$$\hat{\mathbf{w}}^{\text{lasso}} = \operatorname{argmin}_{\mathbf{w}} \left\{ \sum_{i=1:n} (y_i - \mathbf{w}^T \mathbf{x}_i)^2 + \lambda \sum_{j=1:m} |w_j| \right\}$$

Lasso regression (aka L1-regularization)

- Constrains the weights by penalizing the absolute value of their size:

$$\hat{\mathbf{w}}^{lasso} = \operatorname{argmin}_{\mathbf{w}} \{ \sum_{i=1:n} (y_i - \mathbf{w}^T \mathbf{x}_i)^2 + \lambda \sum_{j=1:m} |w_j| \}$$

- Now there is no closed-form solution. Need to solve a quadratic programming problem instead.

Lasso regression (aka L1-regularization)

- Constrains the weights by penalizing the absolute value of their size:

$$\hat{\mathbf{w}}^{\text{lasso}} = \operatorname{argmin}_{\mathbf{w}} \{ \sum_{i=1:n} (y_i - \mathbf{w}^T \mathbf{x}_i)^2 + \lambda \sum_{j=1:m} |w_j| \}$$


- Now there is no closed-form solution. **Need to solve a quadratic programming problem instead.**
 - More computationally expensive than Ridge regression.
 - Effectively sets the weights of less relevant input features to zero.

Comparing Ridge and Lasso

| | Ridge | Lasso |
|--------------------|----------------------|------------------------------|
| <i>Complexity:</i> | $\sum_{j=0:m} w_j^2$ | $\lambda \sum_{j=1:m} w_j $ |

- Note that for lasso, reducing any weight by, say, 0.1 reduces the complexity by the same amount
 - So, we'd prefer reducing less relevant features
- In ridge regression, reducing high weights reduces complexity more than reducing a low weight
 - So, trade-off between reducing less relevant features and larger weights. Tend to not reduce weights that are already small

Comparing Ridge and Lasso



A quick look at evaluation functions

- We call $L(Y, f_w(x))$ the loss function.
 - Least-square / Mean squared-error (MSE) loss:

$$L(Y, f_w(X)) = \sum_{i=1:n} (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

- Other loss functions?

A quick look at evaluation functions

- We call $L(Y, f_w(x))$ the loss function.

- Least-square / Mean squared-error (MSE) loss:

$$L(Y, f_w(X)) = \sum_{i=1:n} (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

- Other loss functions?

- Absolute error loss: $L(Y, f_w(X)) = \sum_{i=1:n} |y_i - \mathbf{w}^T \mathbf{x}_i|$

- 0-1 loss (for classification): $L(Y, f_w(X)) = \sum_{i=1:n} I(y_i \neq f_w(\mathbf{x}_i))$

A quick look at evaluation functions

- We call $L(Y, f_w(X))$ the loss function.

- Least-square / Mean squared-error (MSE) loss:

$$L(Y, f_w(X)) = \sum_{i=1:n} (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

- Other loss functions?

- Absolute error loss: $L(Y, f_w(X)) = \sum_{i=1:n} |y_i - \mathbf{w}^T \mathbf{x}_i|$

- 0-1 loss (for classification): $L(Y, f_w(X)) = \sum_{i=1:n} I(y_i \neq f_w(\mathbf{x}_i))$

- Different loss functions make **different assumptions**.

- Squared error loss assumes the data can be approximated by a global linear model with Gaussian noise.
- Loss function **independent** of complexity penalty (l1 or l2)